

# AASD 4000

# Machine Learning - I

Applied AI Solutions Developer Program



# Module 05

# Scikit Learn

Vejeý Gandýer

# Agenda

Scikit Learn

Importance of Scikit Learn

Setting up scikit-learn

Modeling Process

Estimator API

Use case: Linear Regression problem

Task 6: Linear Regression model

Use case: Unsupervised problem

# Scikit Learn

What is it?



# What is Scikit-learn ?

Scikit-learn (Sklearn) is the most useful and robust library for **machine learning** in Python

It provides a selection of efficient tools for machine learning and statistical modeling including **classification, regression, clustering** and **dimensionality reduction** via a consistence interface in Python

This library, which is largely written in Python, is built upon **NumPy**, and uses **SciPy** and **Matplotlib**.



# Importance of Scikit Learn

Why it is needed?



# Importance of Scikit Learn

## Faster Prototyping

Amazing features enables the machine learning developer to build models **faster**

Main features of scikit-learn are:

- ✓ Dimensionality Reduction
- ✓ Regression
- ✓ Preprocessing
- ✓ Classification
- ✓ Model Selection
- ✓ Clustering

# Setting up Scikit-learn



# Installing Scikit-learn

```
[>>> import sklearn
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
ModuleNotFoundError: No module named 'sklearn'
```

```
[>>> import scikit-learn
      File "<stdin>", line 1
        import scikit-learn
                ^
SyntaxError: invalid syntax
[>>> import sklearn
>>>
```

Conda / Miniconda:

```
conda search scikit-learn
conda install scikit-learn
```

Developer version

```
conda install -c anaconda git
pip install Cython
pip install h5py
pip install git+git://github.com/scikit-learn/scikit-learn.git
```

Virtual environment:

```
pip install -U scikit-learn
```

# Running Scikit learn

```
[In [1]: import scikit-learn
File "<ipython-input-1-56e12c41cf9d>", line 1
import scikit-learn
      ^
SyntaxError: invalid syntax
```

```
[In [2]: import sklearn
```

```
In [3]: █
```

Importing Scikit-learn


Note: sklearn not scikit-learn

# Getting Help

```
[In [4]: from sklearn.]
```

base	covariance	discriminant_analysis	externals	impute
calibration	cross_decomposition	dummy	feature_extraction	isotonic
cluster	datasets	ensemble	feature_selection	kernel_approximation
compose	decomposition	exceptions	gaussian_process	kernel_ridge

# Modeling Process



# Modeling Process

Modeling Process involves the following steps:

1. Load the dataset
2. Preprocess the dataset \*
3. Split the dataset
4. Train the model
5. Persist the model

\* Detailed in a future lecture

# Modeling Process

Loading the dataset



# Loading the dataset

Dataset: Collection of data

Features: Variables of data aka predictors, inputs or attributes

Feature matrix: Collection of features

Feature Names: List of all the names of the features

Response: Output variable depends on feature variables aka target, label or output

Response Vector: Used to represent response column

Target Names: Possible values taken by a response vector

Scikit-learn has in-built datasets

Classification: Iris and digits

Regression: Boston House Prices

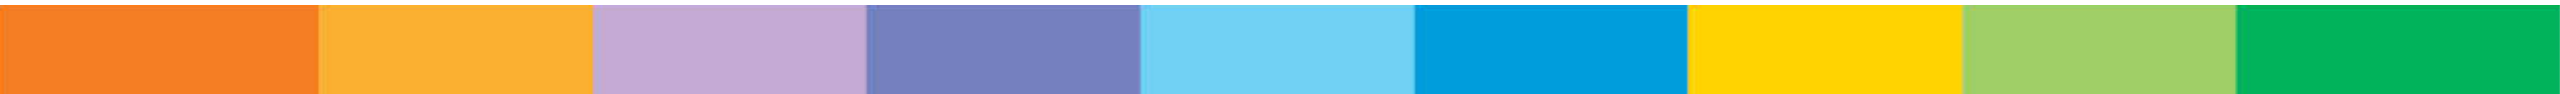
```
from sklearn.datasets import load_iris
iris = load_iris()
X = iris.data
y = iris.target
feature_names = iris.feature_names
target_names = iris.target_names
print("Feature names:", feature_names)
print("Target names:", target_names)
print("\nFirst 10 rows of X:\n", X[:10])
```

```
Feature names: ['sepal length (cm)', 'sepal width (cm)',
'petal length (cm)', 'petal width (cm)']
Target names: ['setosa' 'versicolor' 'virginica']
First 10 rows of X:
[[5.1 3.5 1.4 0.2]
[4.9 3. 1.4 0.2]
[4.7 3.2 1.3 0.2]
[4.6 3.1 1.5 0.2]
[5. 3.6 1.4 0.2]
[5.4 3.9 1.7 0.4]
[4.6 3.4 1.4 0.3]
[5. 3.4 1.5 0.2]
[4.4 2.9 1.4 0.2]
[4.9 3.1 1.5 0.1]]
```




# Modeling Process

Preprocess the dataset



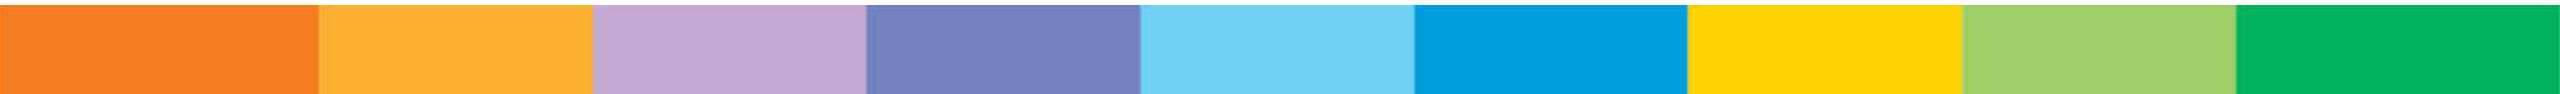
# Data Preprocessing



To be seen in detail in a future lecture

# Modeling Process

Split the dataset



# Splitting the dataset

Training Set: Used to train the model

Testing Set: Used to test the model and not used for training

Split: Ratio of Training Set to Testing Set (70:30 / 80:20)



```
from sklearn.datasets import load_iris
iris = load_iris()
X = iris.data
y = iris.target
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X,
                                                    y,
                                                    test_size=0.3,
                                                    random_state=1
                                                    )
print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(105, 4)
(45, 4)
(105,)
(45,)
```

X

y

test\_size

random\_state

# Modeling Process

Train the model



# Train the model

Model to be trained using the training set

Choose a model to use

Build an object for the model with specific parameters

Fit the model

Predict for the new unseen data



```
# Load dataset
from sklearn.datasets import load_iris
iris = load_iris()
X = iris.data
y = iris.target

# Split dataset
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.4,
random_state=1)

from sklearn.neighbors import KNeighborsClassifier
from sklearn import metrics

# Build the model object
classifier_knn = KNeighborsClassifier(n_neighbors=3)

# Fit the model
classifier_knn.fit(X_train, y_train)

# Predict on testing dataset
y_pred = classifier_knn.predict(X_test)

# Finding accuracy by comparing actual response values(y_test)with predicted
response value(y_pred)
print("Accuracy:", metrics.accuracy_score(y_test, y_pred))

# Providing sample data and the model will make prediction out of that data
sample = [[5, 5, 3, 2], [2, 4, 3, 5]]
preds = classifier_knn.predict(sample)
pred_species = [iris.target_names[p] for p in preds]
print("Predictions:", pred_species)
```

Model - kNN

fit()

predict()

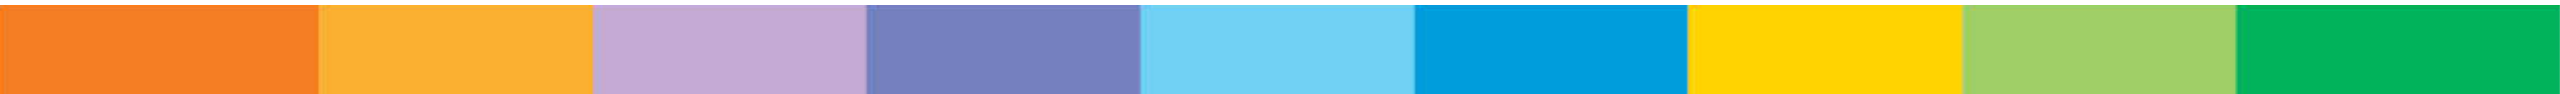
Accuracy: 0.9833333333333333

Predictions: ['versicolor', 'virginica']



# Modeling Process

Persist the model



# Persist the model

Trained Model must be stored and persisted for future use

**joblib** package

**dump()** - save the trained model

**load()** - load the stored model for use

```
# Saving the trained model using dump()
from sklearn.externals import joblib
joblib.dump(classifier_knn, 'iris_classifier_knn.joblib')

# Loading the saved model using load()
joblib.load('iris_classifier_knn.joblib')
```

joblib  
dump()  
load()

# Estimator API

# Steps in using Estimator API

1. Choose a class of model
2. Choose model hyperparameters
3. Arranging the data for the model
4. Model Fitting
5. Applying the model

# #1 - Choosing a model

Choosing a model is as simple as importing the appropriate Estimator class from Sklearn package

Model: **LinearRegression**

```
from sklearn.linear_model import LinearRegression
```

# #2 - Choosing model hyperparameters

Look at all hyperparameters for a model in its signature

```
LinearRegression(  
    fit_intercept=True,  
    normalize=False,  
    copy_X=True,  
    n_jobs=None,  
)
```

Hyperparameter:  
fit\_intercept

## Parameters

-----

**fit\_intercept** : boolean, optional, default True  
whether to calculate the intercept for this model. If set to False, no intercept will be used in calculations (e.g. data is expected to be already centered).

**normalize** : boolean, optional, default False  
This parameter is ignored when ``fit\_intercept`` is set to False. If True, the regressors X will be normalized before regression by subtracting the mean and dividing by the l2-norm. If you wish to standardize, please use :class:`sklearn.preprocessing.StandardScaler` before calling ``fit`` on an estimator with ``normalize=False``.

**copy\_X** : boolean, optional, default True  
If True, X will be copied; else, it may be overwritten.

**n\_jobs** : int or None, optional (default=None)  
The number of jobs to use for the computation. This will only provide speedup for n\_targets > 1 and sufficient large problems. ``None`` means 1 unless in a :obj:`joblib.parallel\_backend` context. ``-1`` means using all processors. See :term:`Glossary <n\_jobs>` for more details.

```
model = LinearRegression(fit_intercept=True)  
model
```

# #3 - Arranging the data

Arrange X, y in its correct dimensions

y - target variable of length n\_samples (1D array)

X - feature matrix of size n\_samples x n\_features

```
X = x[:, np.newaxis]  
X.shape
```

Output:



# #4 - Fitting the model

Model training happens here

fit()

```
model.fit(X, y)
```

Output:

```
model.coef_  
model.intercept_
```

# #5 - Applying the model

Apply the model to the new unseen test data


Output:

Classification problem: `predict()`

Unsupervised problem: `fit()` / `transform()`

```
plt.scatter(x, y)
plt.plot(xfit, yfit)
```

```
xfit = np.linspace(-1, 11)
Xfit = xfit[:, np.newaxis]
yfit = model.predict(Xfit)
```



# Recap: 5 Easy Steps in using Estimator API

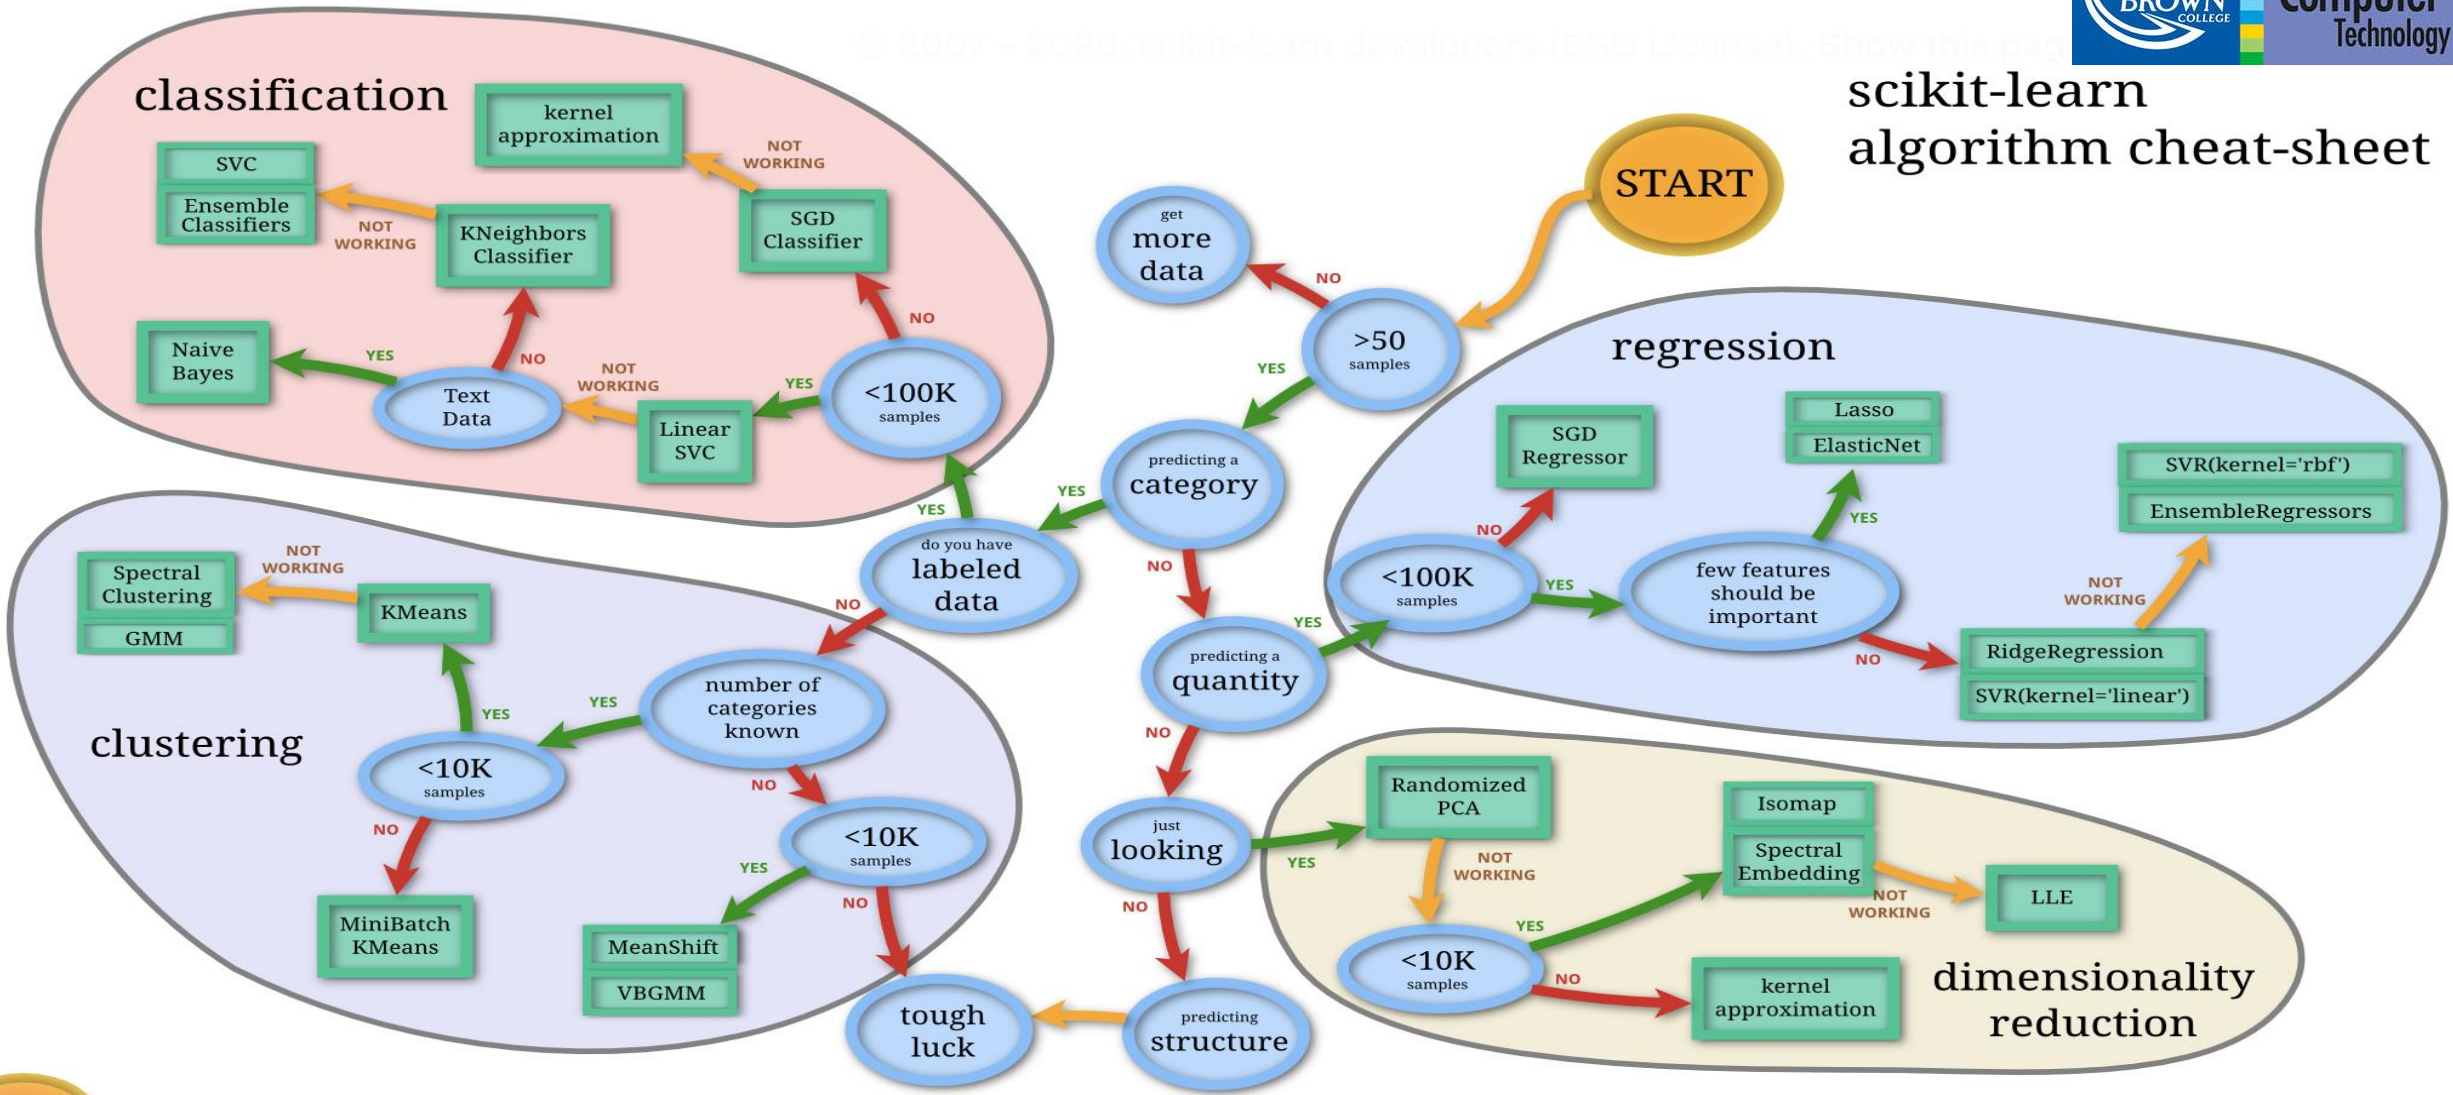
1. Choose a class of model
2. Choose model hyperparameters
3. Arranging the data for the model
4. Model Fitting
5. Applying the model

# Estimator API

Scikit-learn Cheatsheet

[https://scikit-learn.org/stable/tutorial/machine\\_learning\\_map/index.html](https://scikit-learn.org/stable/tutorial/machine_learning_map/index.html)

# scikit-learn algorithm cheat-sheet



# Regression Problem

Task 6: Create a Linear Regression model for Real Estate Housing Prices



# Task 6: Create a Linear Regression model

Use the previous code scripts in creating a Linear Regression model for the below dataset

Real Estate Housing Prices Dataset



# Unsupervised Problem





# Steps in using Estimator API

1. Choose a class of model
2. Choose model hyperparameters
3. Model Fitting
4. Transform the data
5. Visualizing the model

# Choosing a model

Choosing a model is as simple as importing the appropriate Estimator class from Sklearn package

Model: **PCA** (Principal Component Analysis)

```
from sklearn.decomposition import PCA
```

# Choosing model hyperparameters

Look at all hyperparameters for a model in its signature

```
PCA(
    n_components=None,
    copy=True,
    whiten=False,
    svd_solver='auto',
    tol=0.0,
    iterated_power='auto',
    random_state=None,
)
```

Hyperparameter: `n_components`

```
model = PCA(n_components=2)
model
```

## Parameters

`n_components` : int, float, None or string

Number of components to keep.

if `n_components` is not set all components are kept::

```
n_components == min(n_samples, n_features)
```

If ```n_components == 'mle'``` and ```svd_solver == 'full'```, Minka's MLE is used to guess the dimension. Use of ```n_components == 'mle'``` will interpret ```svd_solver == 'auto'``` as ```svd_solver == 'full'```.

If ```0 < n_components < 1``` and ```svd_solver == 'full'```, select the number of components such that the amount of variance that needs to be explained is greater than the percentage specified by `n_components`.

If ```svd_solver == 'arpack'```, the number of components must be strictly less than the minimum of `n_features` and `n_samples`.

Hence, the None case results in::

```
n_components == min(n_samples, n_features) - 1
```

`copy` : bool (default True)

If False, data passed to fit are overwritten and running `fit(X).transform(X)` will not yield the expected results, use `fit_transform(X)` instead.

`whiten` : bool, optional (default False)

When True (False by default) the `'components_'` vectors are multiplied by the square root of `n_samples` and then divided by the singular values to ensure uncorrelated outputs with unit component-wise variances.

Whitening will remove some information from the transformed signal (the relative variance scales of the components) but can sometime improve the predictive accuracy of the downstream estimators by making their data respect some hard-wired assumptions.

`svd_solver` : string {'auto', 'full', 'arpack', 'randomized'}

# Fitting the model

Model training happens here

fit()

```
model.fit(X_iris)
```

Output:

```
PCA(copy=True, iterated_power='auto', n_components=2, random_state=None,  
     svd_solver='auto', tol=0.0, whiten=False)
```

# Transform the data

For PCA, no predict()


Note transform() method

```
X_2D = model.transform(X_iris)
```

# Visualizing the transformed data

```
X_2D = model.transform(X_iris)
```

```
iris['PCA1'] = X_2D[:, 0]  
iris['PCA2'] = X_2D[:, 1]  
sns.lmplot("PCA1", "PCA2",  
           hue='species',  
           data=iris,  
           fit_reg=False  
)
```



# References & Further Reading

## Scikit-learn

[https://scikit-learn.org/stable/user\\_guide.html](https://scikit-learn.org/stable/user_guide.html)

## Kevin Markham (Data School) on using Scikit-learn

<https://www.youtube.com/playlist?list=PL5-da3qGB5lCeMbQuqbbCOQWcS6OYBr5A>